MEGATRON

Large-Scale Language Model Training Framework

TRAINING FRAMEWORK [megatron/training/]

- Training Loop — *training.py*
- Arguments — *arguments.py*
- Checkpointing — *checkpointing.py*
- Initialization — *initialize.py*

MODELS

MCore

- GPT — *core/models/gpt/*
- BERT — *core/models/bert/*
- T5 — *core/models/t5/*
- Mamba — *core/models/mamba/*
- Multimodal — *core/models/multimodal/*
- MoE — *core/models/mimo/*

TRANSFORMER CORE

MCore

- **Attention** — *core/transformer/attention.py* — MCore
- **MLP** — *core/transformer/mlp.py* — MCore
- **Layer Norm** — *core/transformer/* — MCore
- **Embeddings** — *core/transformer/* — MCore
- **Block** — *core/transformer/* — MCore

PARALLELISM STRATEGY [megatron/core/]

TENSOR PARALLELISM

Intra-layer model sharding

core/tensor_parallel/ | layers.py | mappings.py

PIPELINE PARALLELISM

Inter-layer distribution

core/pipeline_parallel/ | schedules.py

DATA PARALLELISM

Gradient synchronization

core/distributed/ | ddp.py | fsdp.py

DATA PIPELINE

MCore

- Blended Dataset — *core/datasets/blended_* — MCore
- Indexed Dataset — *core/datasets/indexed_* — MCore
- GPT Dataset — *core/datasets/gpt_* — MCore
- Data Loader — *training/datasets/*

OPTIMIZER

MCore

- Distributed Optimizer — *core/optimizer/* — MCore
- Adam Config — *core/optimizer/* — MCore
- Gradient Scaling — *core/optimizer/* — MCore

CHECKPOINTING [training/]

- Save/Load — *training/*
- Distributed — *core/dist_checkpointing/* — MCore
- Model State — *training/*

KEY FEATURES [megatron/core/]

PRECISION SUPPORT

MCore

- FP16 / BF16 Training
- FP8 (Hopper Optimized)
- FP4 Quantization
- Transformer Engine

MEMORY OPTIMIZATION

MCore

- Activation Checkpointing
- Sequence Parallelism
- Gradient Checkpointing
- Fine-grained Offloading

PERFORMANCE

MCore

- CUDA Graphs
- Fused Kernels
- Flash Attention
- Fused LayerNorm

Megatron: Large-Scale Language Model Training Framework
NVIDIA Research | Distributed Training | Tensor | Pipeline | Data Parallelism

ARCHITECTURAL PATTERNS

Modular Core Library

MCore

megatron/core/ provides production-ready, GPU-optimized building blocks for framework developers

models | transformer | tensor_parallel
pipeline_parallel | distributed | optimizer

Configuration-Driven

Dataclass-based configuration system enables flexible model parallelism setup via command-line arguments

TransformerConfig | ModelParallelConfig
ProcessGroupCollection | parallel_state